Featured

Getting started

Hello, world

Simple web scraper

Serving web endpoints

Large language models (LLMs)

Featured examples

Serving Diffusion models

Serve Stable Diffusion XL on Modal with a number of optimizations for blazingly fast inference.

View on GitHub

Serverless TensorRT-LLM

Run large language models at SOTA performance by building a TRT-LLM engine on Modal.

View on GitHub

Stable Diffusion fine-tuning with Dreambooth

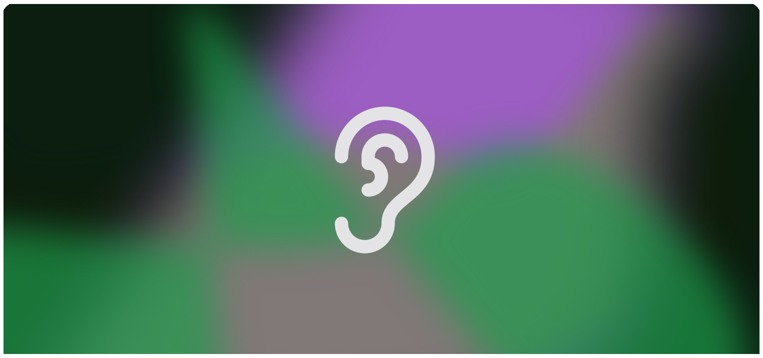

Voice chat with LLMs

Build a real-time voice chat app by combining speech-to-text, an LLM, and text-to-speech.

Fine-tune a Stable Diffusion model on images of your pet using Dreambooth.

View on GitHub

View on GitHub

Fast podcast transcriptions

Build an end-to-end podcast transcription app that leverages dozens of containers for super-fast processing.

View on GitHub

Document OCR job queue

Use Modal as an infinitely scalable job queue that can service async tasks from a web app.

View on GitHub

Parallel processing of Parquet files on S3

Analyze data from the Taxi and Limousine Commission of NYC in parallel.

View on GitHub

Retrieval-Augmented Generation for Q&A

Build a question-answering web endpoint that can cite its sources.

View on GitHub

Hacker News Slackbot

Real-time Object Detection

Use Modal to deploy a cron job that periodically queries Hacker News for new posts, and posts the results to Slack.

View on GitHub

Create a web endpoint that leverages HuggingFace models to do object detection in real-time.

View on GitHub

ControlNet playgrounds

Play with all 10 demo Gradio apps from the ControlNet project.

View on GitHub

About

Slack Community

Status

Pricing

Changelog

Examples

Documentation